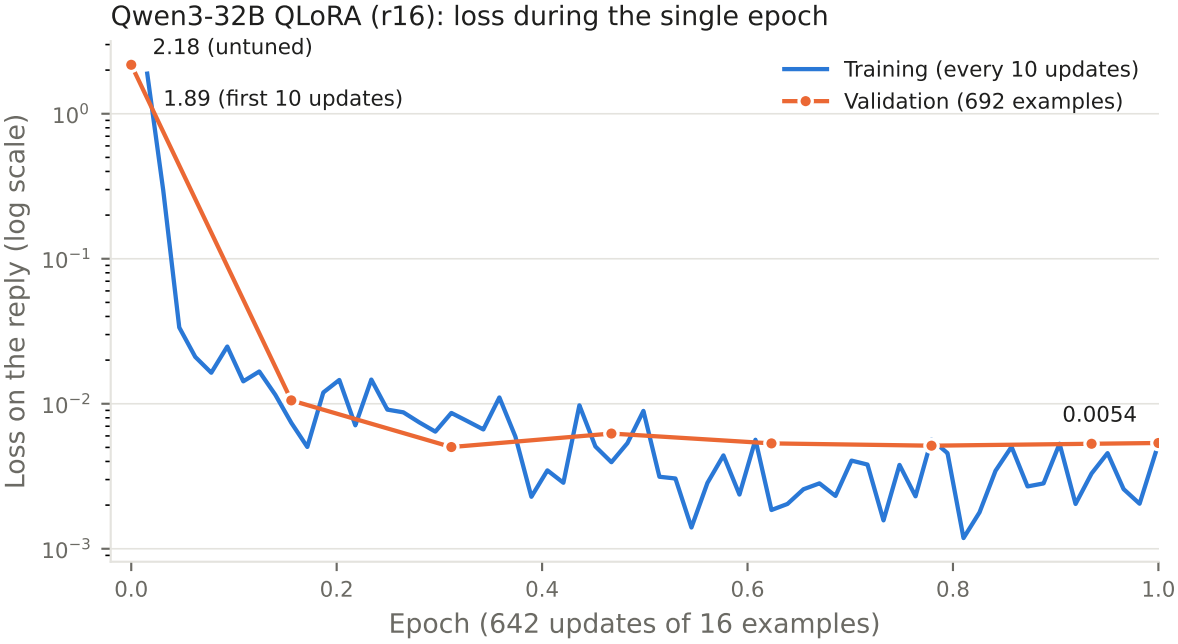
Qwen3-32B QLoRA (r16): loss during the single epoch
2.18 (untuned)
1.89 (first 10 updates)
Training (every 10 updates)
Validation (692 examples)
0.0054
Loss on the reply (log scale)
10^{0}
10^{-1}
10^{-2}
10^{-3}
0.0
0.2
0.4
0.6
0.8
1.0
Epoch (642 updates of 16 examples)